ECE408 Final Project Report

Shriyaa Mittal (smittal6), Tiancheng Wu (twu54), Chuankai Zhao (czhao37)

October 27, 2018

Team Name: smartconvolutionteam

Milestone 1: Due October 24, 2018

Include a list of all kernels that collectively consume more than 90% of the program time

Top 10 kernels are as below:

1. volta_scudnn_128x32_relu_interior_nn_v1

2. Implicit_convolve_sgemm

3. volta_sgemm_128x128_tn

4. activation_fw_4d_kernel

5. pooling_fw_4d_kernel

6. MapPlanLargeKernel

7. SoftmaxKernel

8. MapPlanKernel

9. volta_sgemm_32x32_sliced1x4_tn

10. computeOffsetsKernel

Include a list of all CUDA API calls that collectively consume more than 90% of the program time

Top 10 CUDA API calls are as below:

1. cudaStreamCreateWithFlags

2. cudaMemGetInfo

3. cudaFree

4. cudaEventCreateWithFlags

5. cudaMemcpy2DAsync

6. cudaFuncSetAttribute

7. cudaStreamSynchronize

8. cudaMalloc

9. cudaGetDeviceProperties

10. cudaMemcpy

Include an explanation of the difference between kernels and API calls

Kernels are programmer defined functions, while API calls are built-in.

Show output of rai running MXNet on the CPU

```
* Running /usr/bin/time python m1.1.py
Loading fashion-mnist data... done
Loading model... done
New Inference
EvalMetric: {'accuracy': 0.8177}
```

List program run time

```
19.48user 4.09system 0:13.30elapsed 177%CPU
```

Show output of rai running MXNet on the GPU

```
* Running /usr/bin/time python m1.2.py
Loading fashion-mnist data... done
Loading model... done
New Inference
EvalMetric: {'accuracy': 0.8177}
```

List program run time

```
4.05user 2.67system 0:04.63elapsed 145%CPU
```

MileStone 2: Due Octber 29, 2018

List whole program execution time

```
157.56user 4.42system 2:31.72elapsed 106%CPU
```

List Op times

```
Op Time: 28.737308
Op Time: 118.727850
Correctness: 0.8171 Model: ece408
```